

Medgema Documentation

Objective audience:

- Client non technical.

Index:

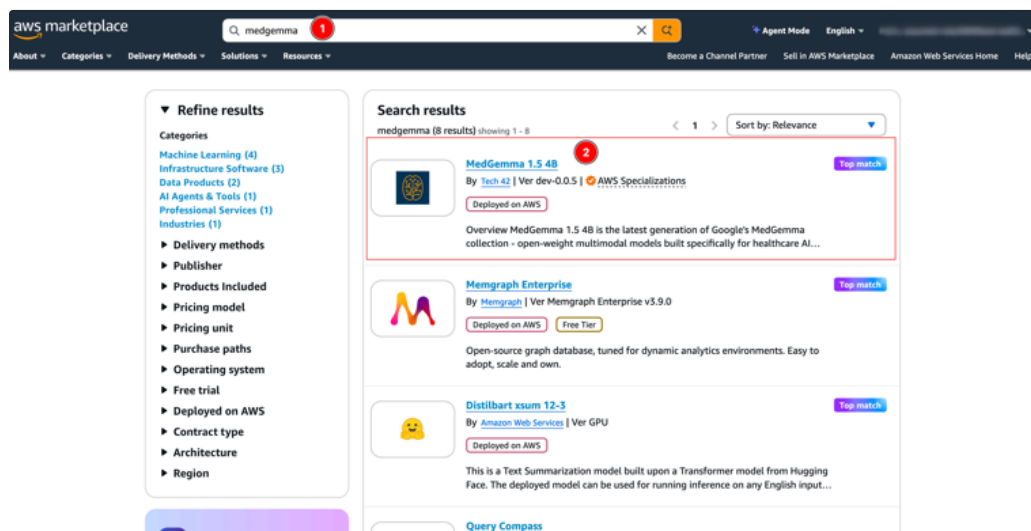
1. Pre-requisites
 - a. Quota for Sagemaker RT or Batch transform job.
2. Subscribe to the MedGemma in marketplace
3. Deploy the product
4. Known Issues
 - a. Quota not available
 - b. Insufficient quota in the region.

Subscribe to the MedGemma marketplace product

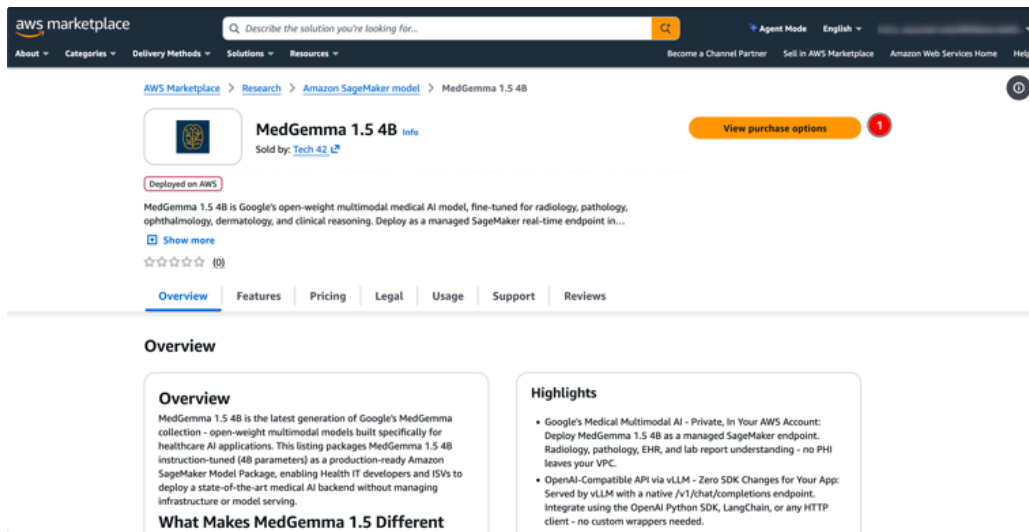
This guide walks you through subscribing to MedGemma via AWS Marketplace, deploying it to a SageMaker endpoint, and running medical image inference — all from your existing AWS account.

Step 1:

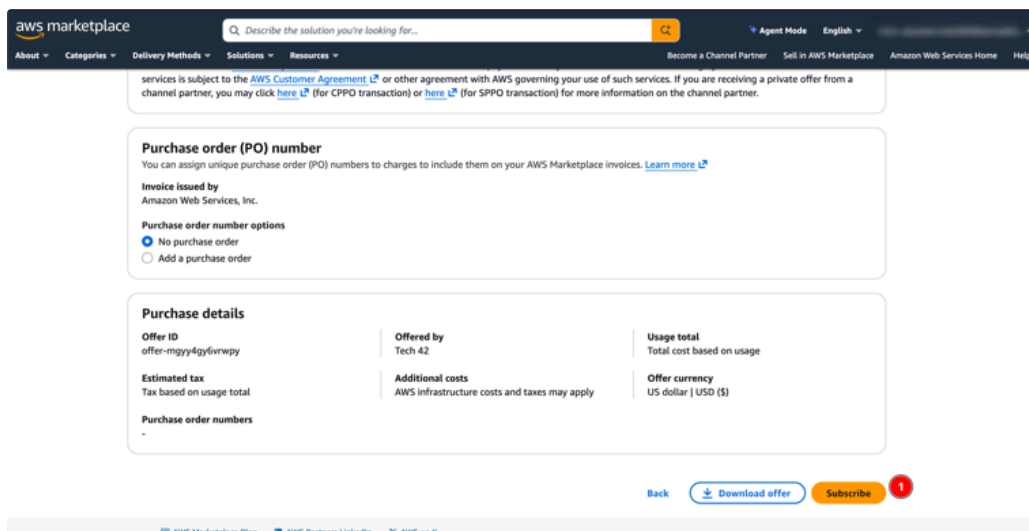
1. Go to [AWS Marketplace](https://aws.amazon.com/marketplace) (<https://aws.amazon.com/marketplace>)
2. Search for "**MedGemma**" or navigate directly to the model listing.
3. Choose the Medgemma product offered by Tech42.



Step 2: On the listing page, click **"View purchase options."**



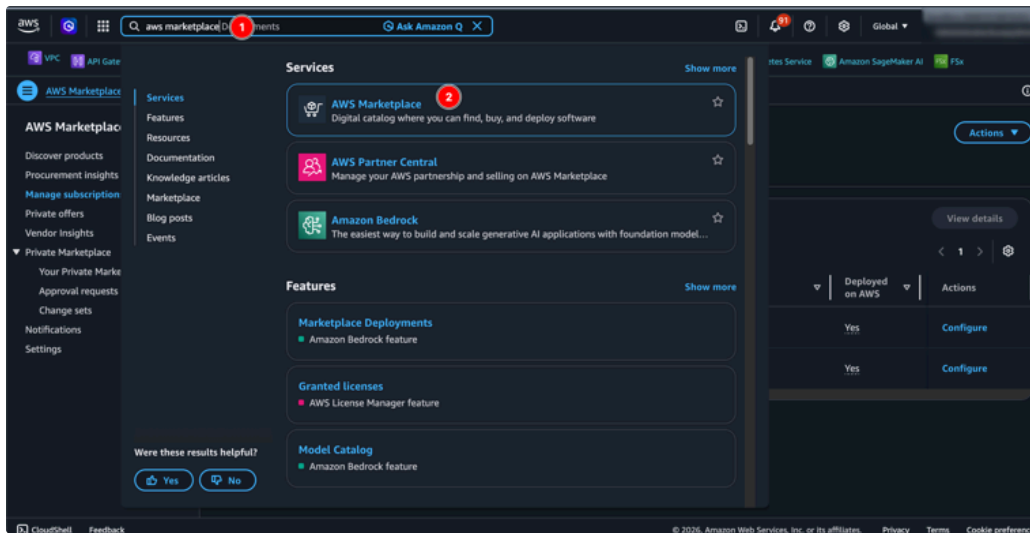
Step 3: Review the End User License Agreement (EULA) and pricing. Click **"Subscribe"** to proceed.



Wait 1–2 minutes for the subscription to validate in your account.

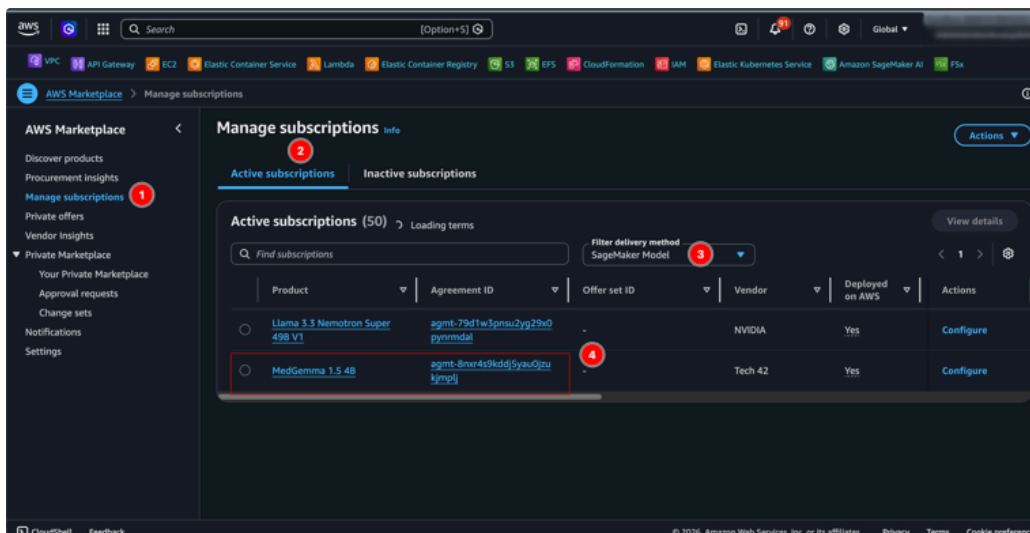
Step 4: Once subscribed, go to

1. **Type AWS Marketplace**
2. Select the AWS Marketplace option



Step 5: Go to

1. **Manage Subscriptions**
2. **Active Subscriptions**
3. **Filter delivery method by “Sagemaker Model”**
4. Confirm MedGemma appears in your product list.



Deploying MedGemma

Get Model ARN

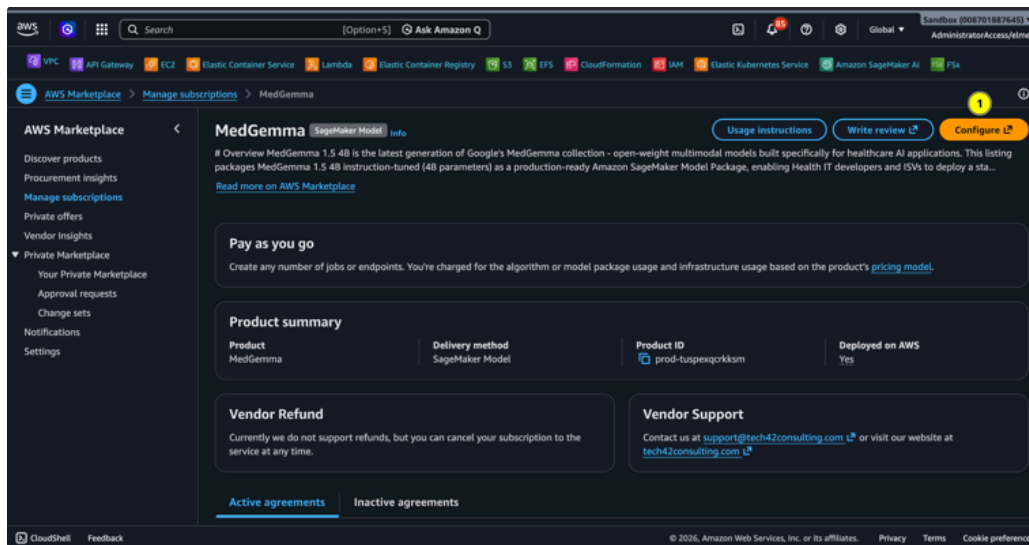
Deploy CF

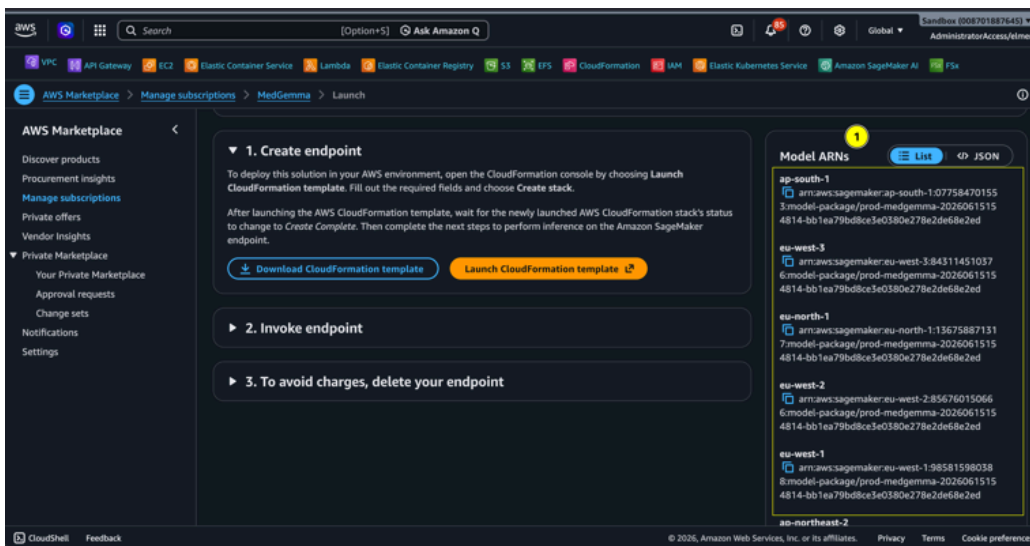
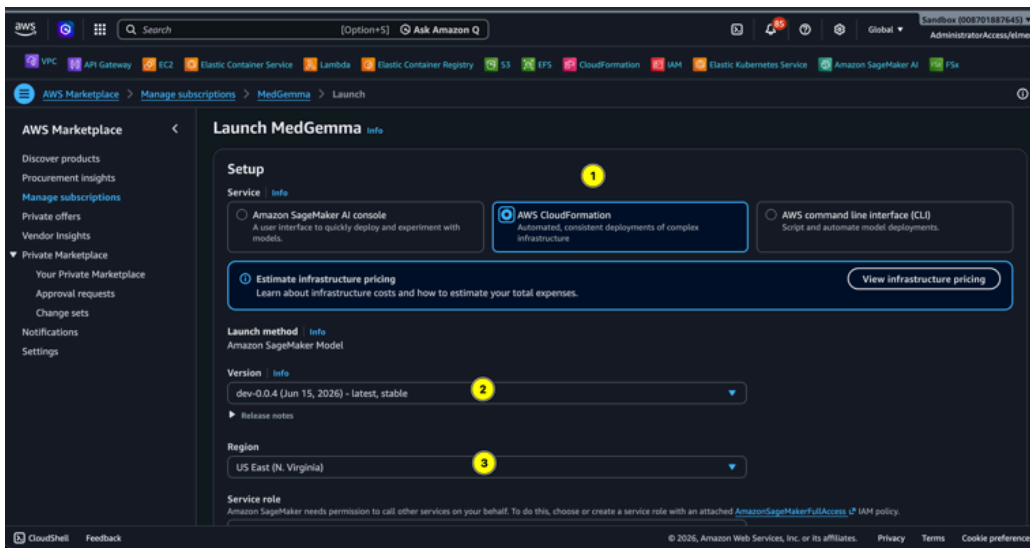
Use the notebook

Get the Model ARN (Pre-requisite)

Before proceeding with the product deployment, we would need to get the product ARN. This will change depending on the region choose to deploy it, and the version of the product that we are deploying. To get the Model ARN, please go to the

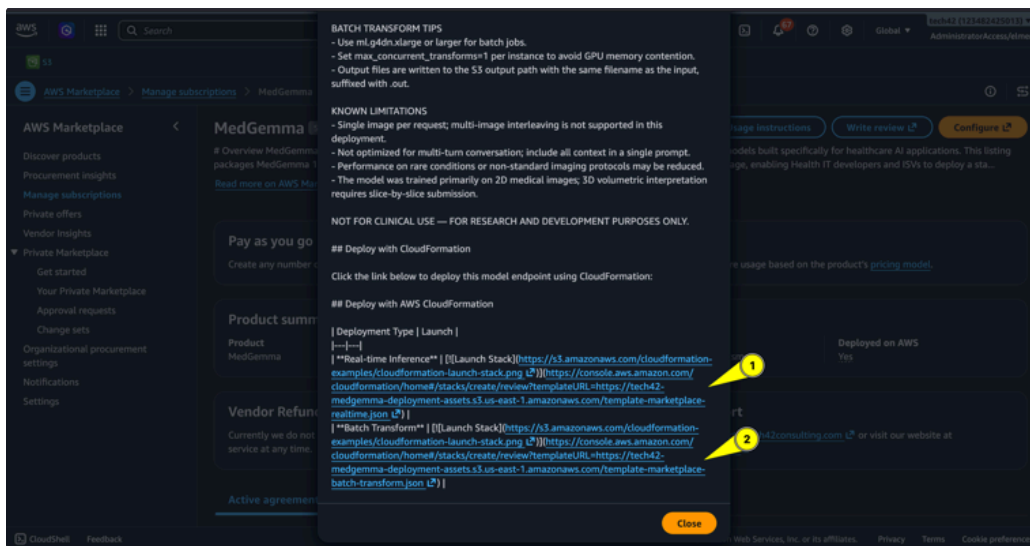
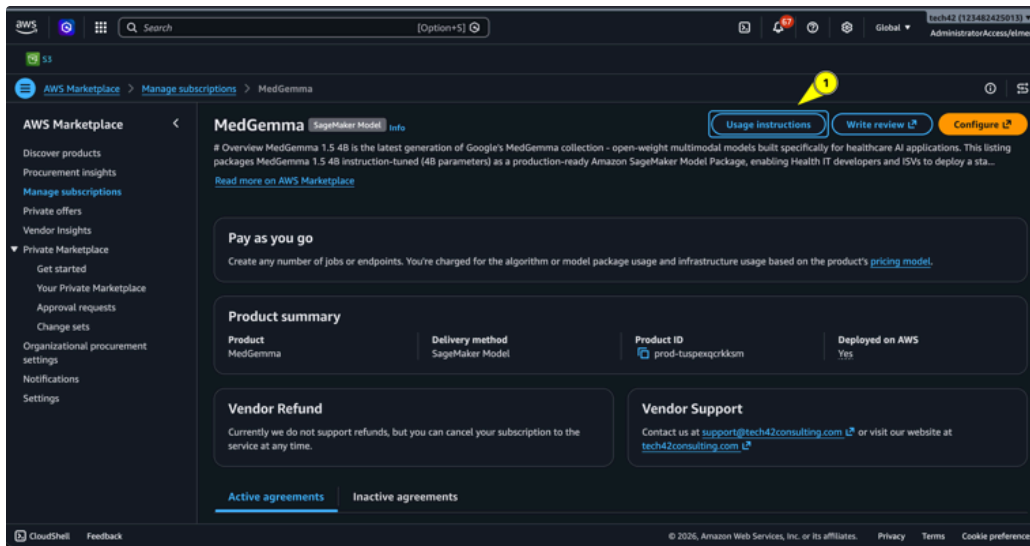
Configure → AWS CloudFormation → Choose the version → Choose the region → Scroll down to the Model ARNs and copy the one that for the region where you are going to deploy the solution





Deploy the Endpoint

To start deploying the solution, you can go to **Usage instructions** → It will pop-up a menu, scroll down, and you will find the link for the cloudformation templates for the real time (1) and batch transform jobs (2) endpoints.



Deploy Real Time Endpoint

After clicking the real time endpoint link

```

1
2 | Deployment Type | Launch |
3 |---|---|
4 | **Real-time Inference** | [[Launch Stack]
  | (https://s3.amazonaws.com/cloudformation-examples/cloudformation-
  | launch-stack.png)]
  | (https://console.aws.amazon.com/cloudformation/home#/stacks/create/review?templateURL=https://tech42-medgemma-deployment-assets.s3.us-east-1.amazonaws.com/template-marketplace-realtime.json) |
5 | **Batch Transform** | [[Launch Stack]
  | (https://s3.amazonaws.com/cloudformation-examples/cloudformation-
  | launch-stack.png)]
  | (https://console.aws.amazon.com/cloudformation/home#/stacks/create/review?templateURL=https://tech42-medgemma-deployment-assets.s3.us-east-1.amazonaws.com/template-marketplace-batch-transform.json) |

```

Set **Stack Name** (1)

Quick create stack

Template

Template URL
https://tech42-medgemma-deployment-assets.s3.us-east-1.amazonaws.com/template-marketplace-realtime.json

Stack description
Creates a SageMaker real-time endpoint from an AWS Marketplace model package. Use this template after subscribing to the product in AWS Marketplace.

► View template

Provide a stack name

Stack name
test-medgemma 1

Stack name can include letters (A-Z and a-z), numbers (0-9), and hyphens.

Parameters

Parameters are defined in your template and allow you to input custom values when you create or update a stack.

ContainerStartupHealthCheckTimeoutInSeconds
Maximum time SageMaker waits for the container to pass startup health checks.
1800

Fill the Parameters:

Parameters

Parameters are defined in your template and allow you to input custom values when you create or update a stack.

ContainerStartupHealthCheckTimeoutInSeconds
Maximum time SageMaker waits for the container to pass startup health checks.
1800

EnableAutoScaling
Whether to configure Application Auto Scaling for the endpoint variant.
Yes

EndPointName
Name for the SageMaker endpoint.
medgemma-marketplace-endpoint

IAMRole
Optional SageMaker execution role ARN. Leave blank to let this stack create one.
Enter String

InstanceCount
Initial instance count for the endpoint.
1

InstanceType
SageMaker real-time inference instance type.
ml.g7e.2xlarge

Add the **ProductARN** copied in the pre-requisite step.

InvocationsPerInstanceTarget
Target invocations per instance for real-time endpoint autoscaling.
5

MaxInstanceCount
Maximum instance count for endpoint autoscaling.
4

MinInstanceCount
Minimum instance count for endpoint autoscaling.
1

ModelDataDownloadTimeoutInSeconds
Maximum time SageMaker waits for model package artifacts to download.
3600

ProductARN
AWS Marketplace product/model package ARN for the subscribed software version and region. Copy this from the AWS Marketplace configuration page.
arn:aws:sagemaker:us-west-1:382657785993:model-package/prod-medgemma-20260615154814-bb1ea79bd8ce3e0380e278e2de8e2ed 1

Referrer
Optional link back to the Marketplace product configuration page.
Enter String

ScaleInCooldownInSeconds
Cooldown, in seconds, after a scale-in activity.
300

Go to Create Stack

